


# Chapter 3

## Transport Layer

Yaxiong Xie

Department of Computer Science and Engineering  
University at Buffalo, SUNY

Adapted from the slides of the book's authors

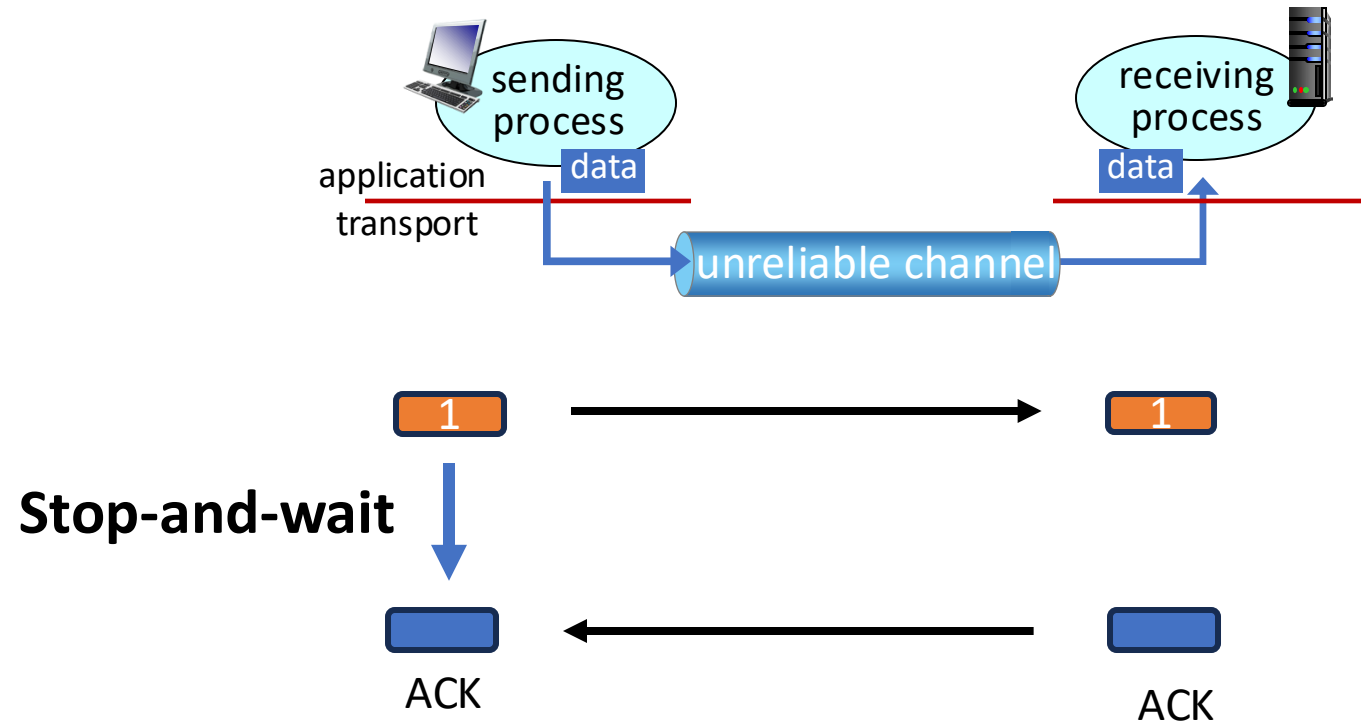


*Computer Networking: A  
Top-Down Approach*


8<sup>th</sup> edition

Jim Kurose, Keith Ross  
Pearson, 2020

# Stop-and-wait operation




# Stop-and-wait operation



# Pipelined protocols operation


**pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged packets

- range of sequence numbers must be increased
- buffering at sender and/or receiver




(a) a stop-and-wait protocol in operation

# Pipelining



How many packets shall we send before receiving the ACK?




# Sliding Window



# Sliding Window



# Sliding Window






# Sliding Window




# Sliding Window




- Timer for oldest in-flight packet
- $Timeout(n)$ : retransmit packet  $n$  and all higher seq # packets in window


# Sliding Window + cumulative ACK




# Sliding Window + cumulative ACK



# Sliding Window + cumulative ACK




# Sliding Window + cumulative ACK



- **cumulative ACK:**  $ACK(n)$ : ACKs all packets up to, including seq #  $n$ 
  - on receiving  $ACK(n)$ : move window forward to begin at  $n+1$


# Sliding Window + cumulative ACK: out-of-order



## Two Question:

1. do we need to send ACK when we receive an out-of-order packet? If yes, what kind of ACK?

# Sliding Window + cumulative ACK: out-of-order




## Two Question:

1. do we need to send ACK when we receive an out-of-order packet? If yes, what kind of ACK?
2. how shall we deal with the out-of-order packets we received?




# Sliding Window + cumulative ACK: out-of-order



## Two Question:


1. do we need to send ACK when we receive an out-of-order packet? If yes, what kind of ACK?
2. how shall we deal with the out-of-order packets we received?

# Sliding Window + cumulative ACK: **out-of-order**



- ACK-only: always send ACK for correctly-received packet so far, with highest **in-order** seq #
  - may generate duplicate ACKs


# Sliding Window + cumulative ACK: out-of-order



- on receipt of out-of-order packet:
  - can discard (don't buffer) or buffer: an implementation decision
  - re-ACK pkt with highest in-order seq #

# Go-Back-N: sender

- sender: “window” of up to  $N$ , consecutive transmitted but unACKed pkts
  - $k$ -bit seq # in pkt header

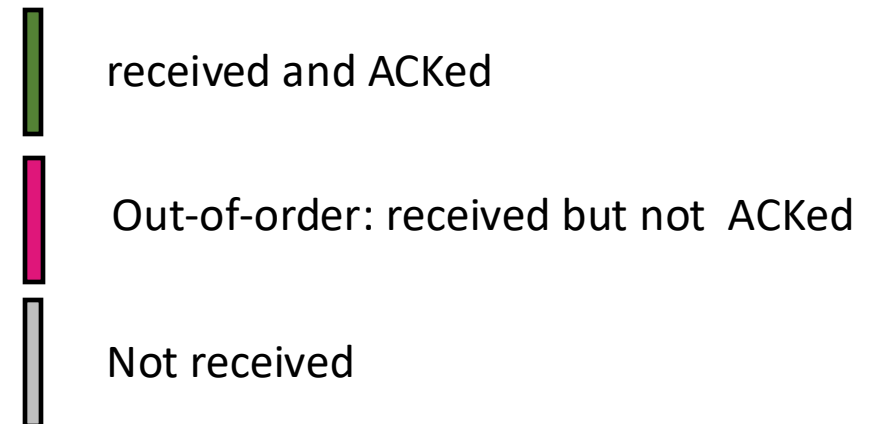



- ***cumulative ACK***:  $ACK(n)$ : ACKs all packets up to, including seq #  $n$ 
  - on receiving  $ACK(n)$ : move window forward to begin at  $n+1$
- timer for oldest in-flight packet
- *timeout*( $n$ ): retransmit packet  $n$  and all higher seq # packets in window

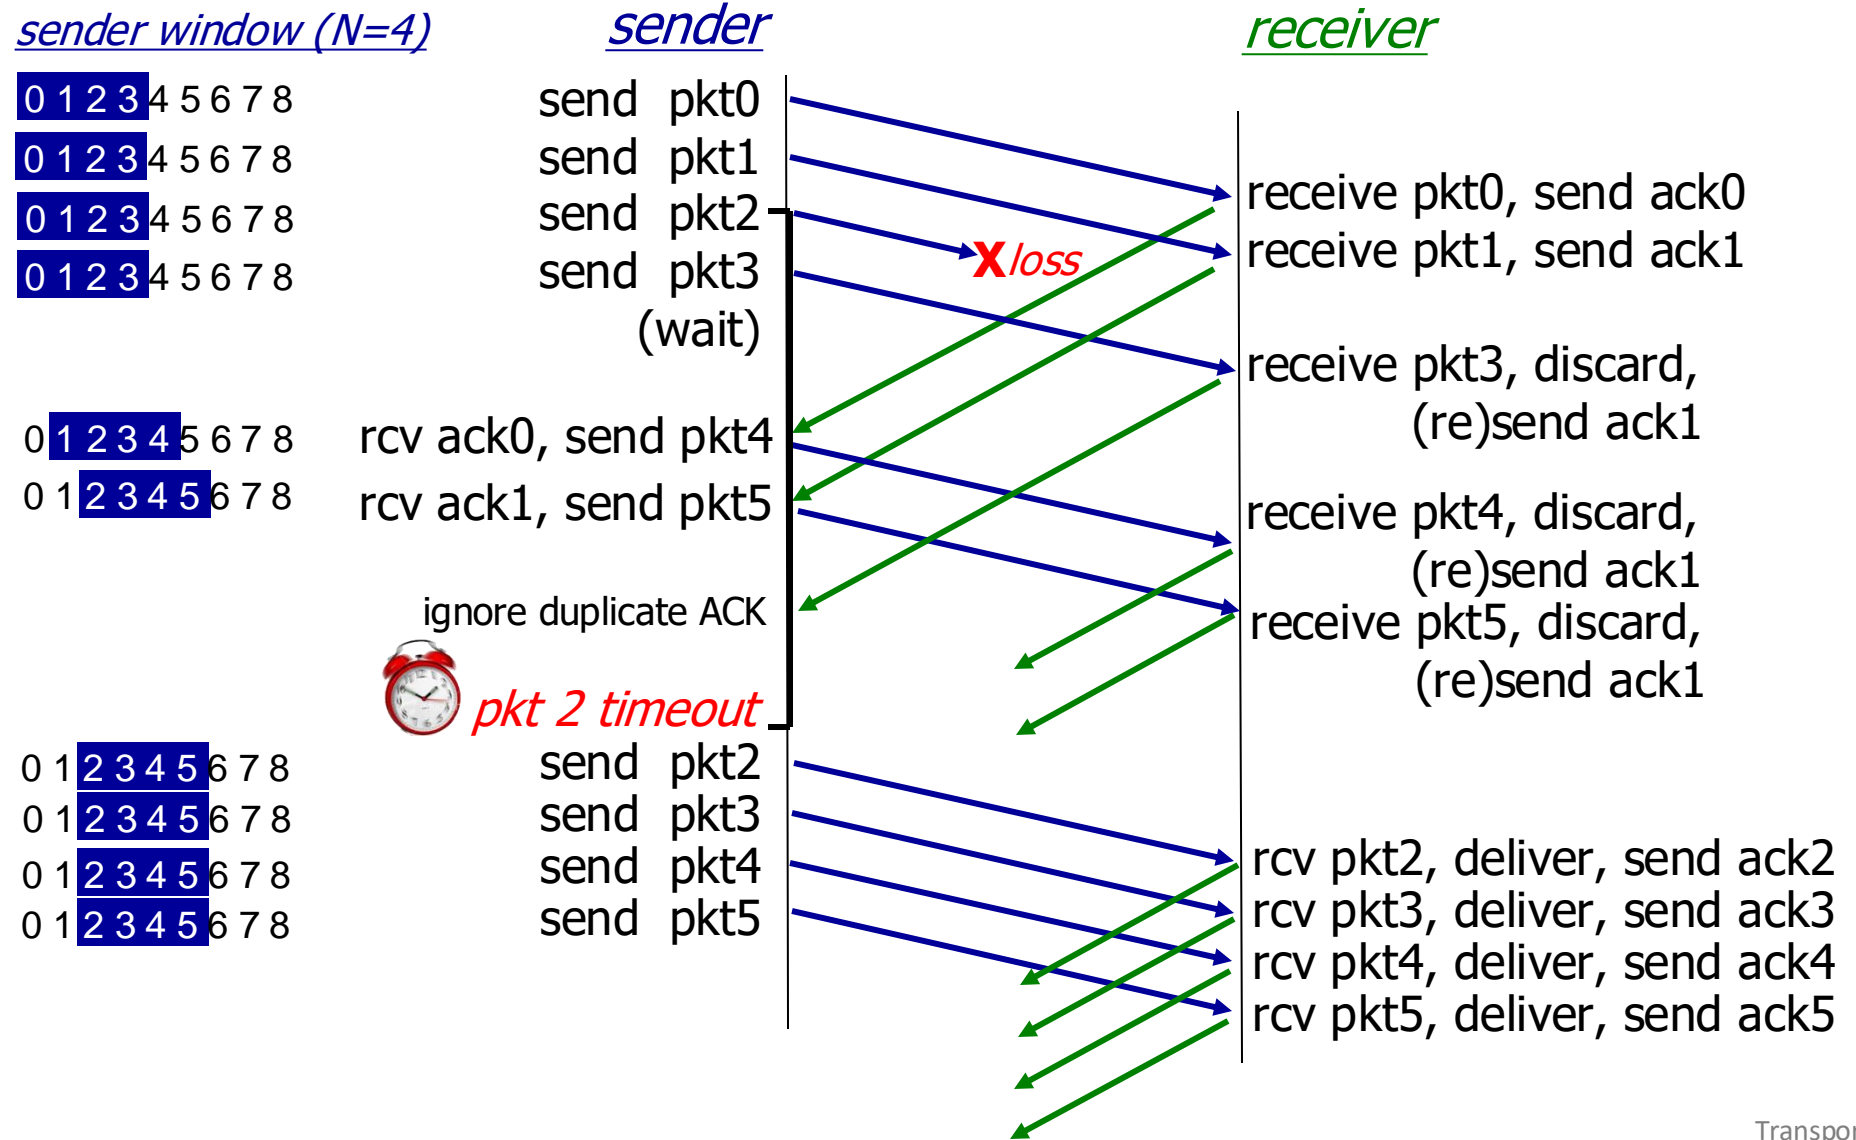
# Go-Back-N: receiver

- ACK-only: always send ACK for correctly-received packet so far, with highest *in-order* seq #
  - may generate duplicate ACKs
  - need only remember `rcv_base`
- on receipt of out-of-order packet:
  - can discard (don't buffer) or buffer: an implementation decision
  - re-ACK pkt with highest in-order seq #

Receiver view of sequence number space:




# Go-Back-N in action




# Sliding Window + selective repeat

# Sliding Window + selective repeat







# Sliding Window + selective repeat




# Sliding Window + selective repeat




# Sliding Window + selective repeat: out-of-order



# Sliding Window + selective repeat: out-of-order




# Sliding Window + selective repeat: **out-of-order**




- ***receiver individually ACKs*** all correctly received packets
  - buffers packets, as needed, for in-order delivery to upper layer

# Sliding Window + selective repeat: **out-of-order**




- ***receiver individually ACKs*** all correctly received packets
  - buffers packets, as needed, for in-order delivery to upper layer

# Sliding Window + selective repeat: **out-of-order**



- ***receiver individually ACKs*** all correctly received packets
  - buffers packets, as needed, for in-order delivery to upper layer

# Sliding Window + selective repeat: **out-of-order**



- ***receiver individually ACKs*** all correctly received packets
  - buffers packets, as needed, for in-order delivery to upper layer




# Sliding Window + selective repeat: **out-of-order**




- ***receiver individually ACKs*** all correctly received packets
  - buffers packets, as needed, for in-order delivery to upper layer

# Sliding Window + selective repeat: **out-of-order**




- ***receiver individually ACKs*** all correctly received packets
  - buffers packets, as needed, for in-order delivery to upper layer

# Sliding Window + selective repeat: **out-of-order**



- ***receiver individually ACKs*** all correctly received packets
  - buffers packets, as needed, for in-order delivery to upper layer

# Sliding Window + selective repeat




- Maintains (conceptually) a timer for each unACKed pkt
  - timeout: retransmits single unACKed packet associated with timeout

# Selective repeat: the approach

- *pipelining*: multiple packets in flight
- *receiver individually ACKs* all correctly received packets
  - buffers packets, as needed, for in-order delivery to upper layer
- sender:
  - maintains (conceptually) a timer for each unACKed pkt
    - timeout: retransmits single unACKed packet associated with timeout
  - maintains (conceptually) “window” over *N* consecutive seq #s
    - limits pipelined, “in flight” packets to be within this window

# Selective repeat: sender, receiver windows



# Selective repeat: sender and receiver

## sender

### data from above:

- if next available seq # in window, send packet

### timeout( $n$ ):

- resend packet  $n$ , restart timer

### ACK( $n$ ) in [sendbase, sendbase+N-1]:

- mark packet  $n$  as received
- if  $n$  smallest unACKed packet, advance window base to next unACKed seq #

## receiver

### packet $n$ in [rcvbase, rcvbase+N-1]

- send ACK( $n$ )
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order packets), advance window to next not-yet-received packet

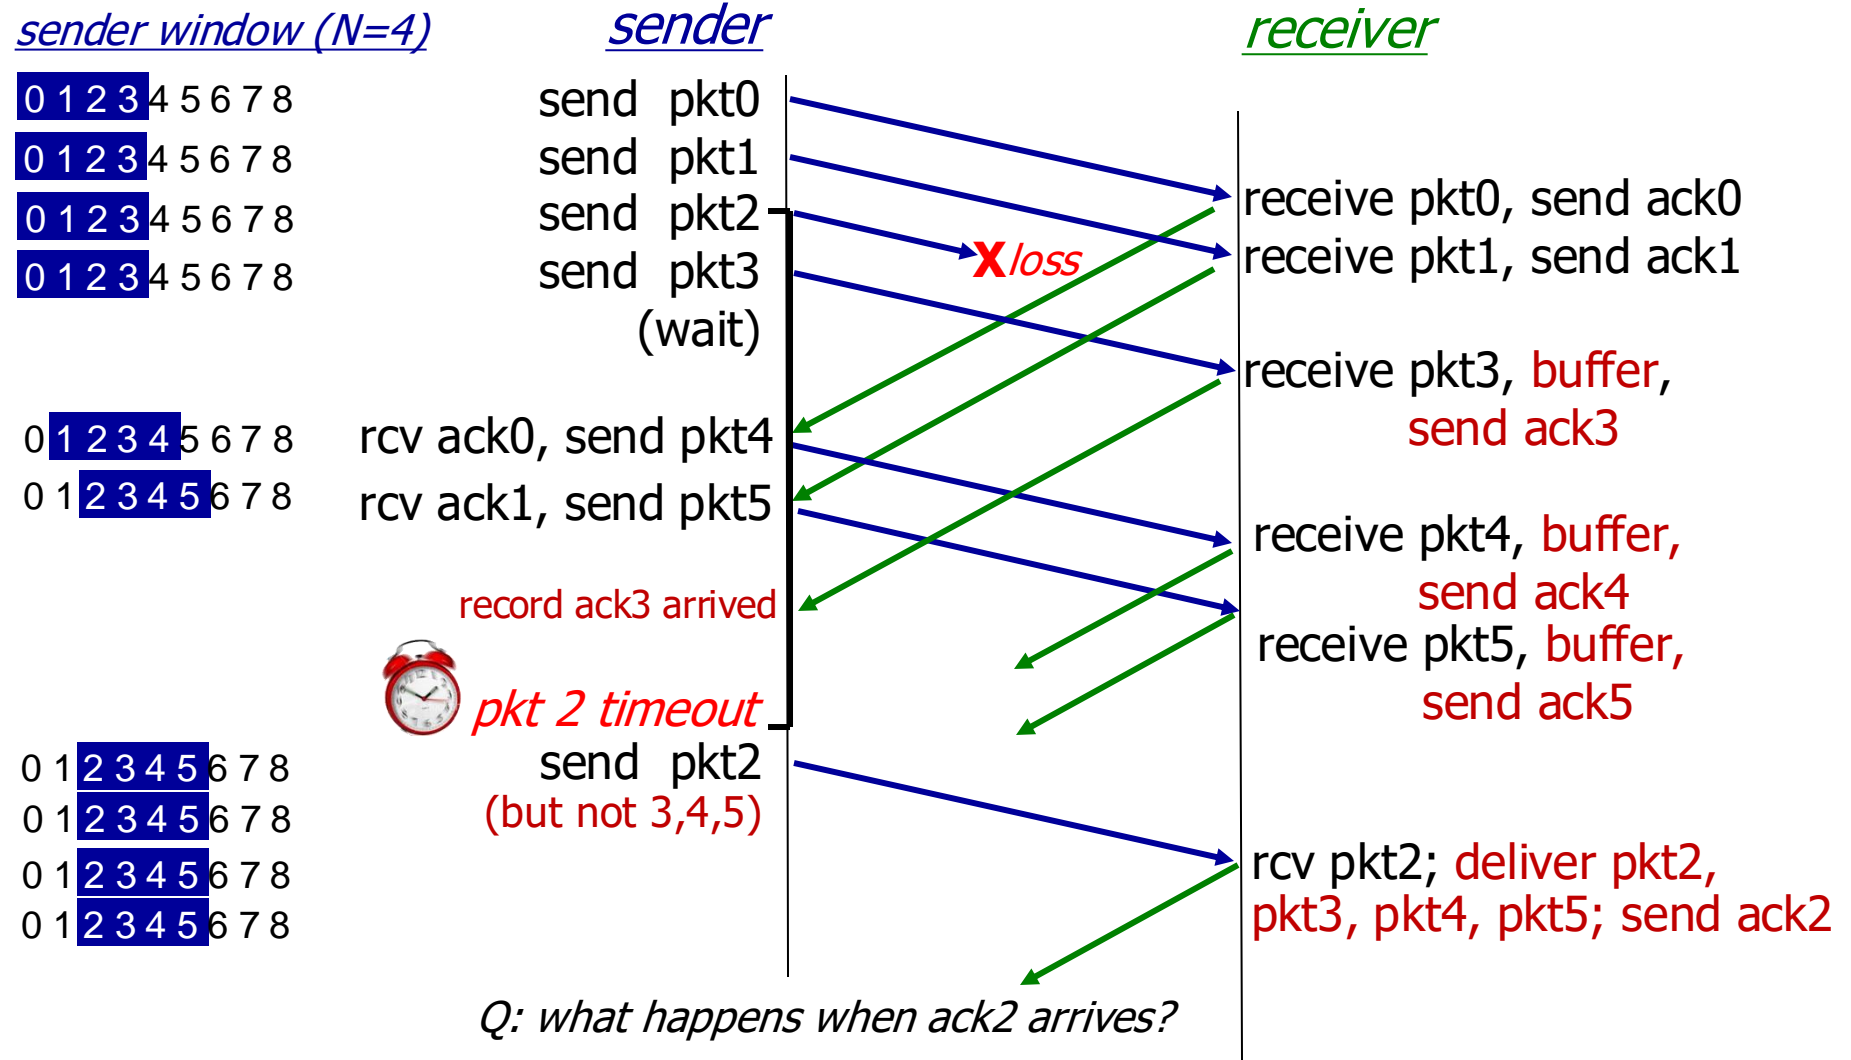
### packet $n$ in [rcvbase-N, rcvbase-1]

- ACK( $n$ )

### otherwise:

- ignore

# Selective Repeat in action

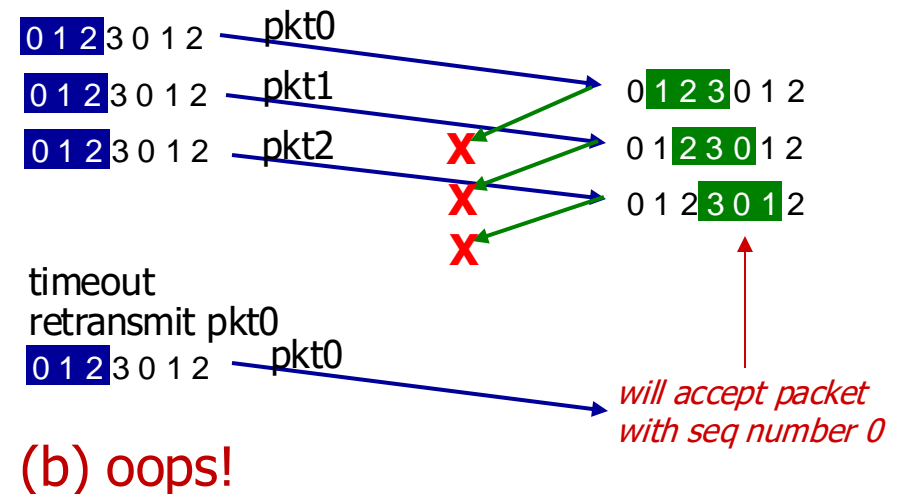
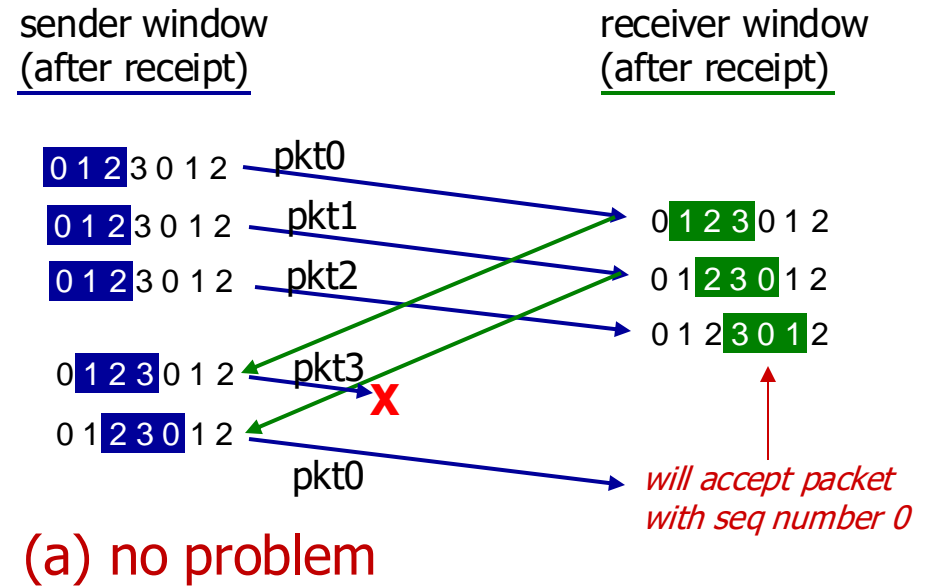




# Selective repeat: a dilemma!

example:

- seq #s: 0, 1, 2, 3 (base 4 counting)
- window size=3




# Selective repeat: a dilemma!

example:


- seq #s: 0, 1, 2, 3 (base 4 counting)
- window size=3

sender window  
(after receipt)

receiver window  
(after receipt)



- receiver can't see sender side
- receiver behavior identical in both cases!
- *something's (very) wrong!*



(b) oops!

# Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- **Connection-oriented transport: TCP**
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- Principles of congestion control
- TCP congestion control

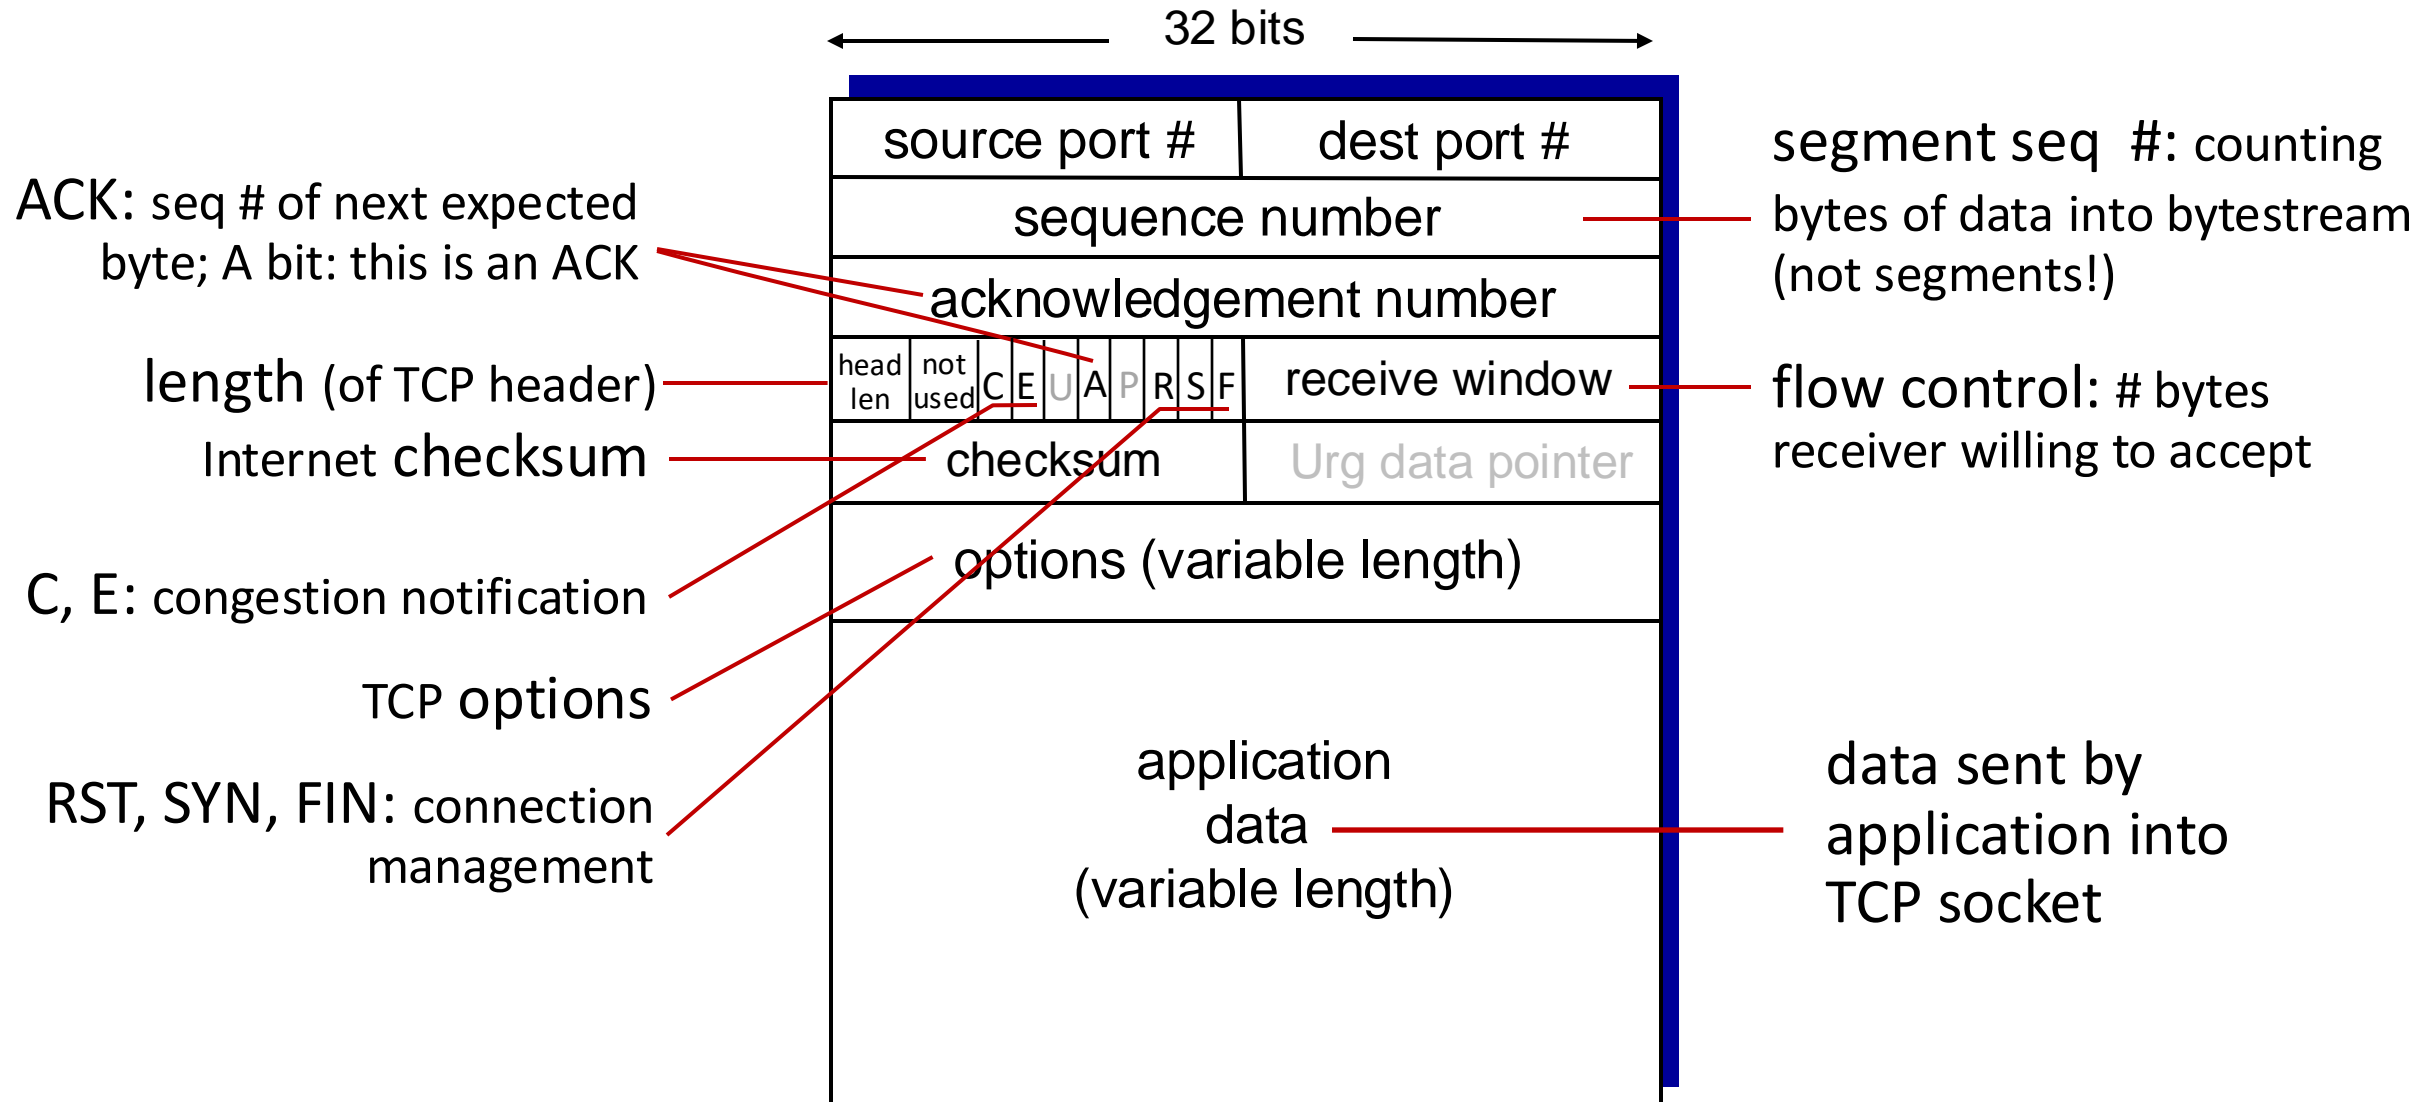


# TCP: overview

RFCs: 793, 1122, 2018, 5681, 7323

- **point-to-point:**
  - one sender, one receiver
- **reliable, in-order *byte stream*:**
  - no “message boundaries”
- **full duplex data:**
  - bi-directional data flow in same connection
  - MSS: maximum segment size
- **cumulative ACKs**
- **pipelining:**
  - TCP congestion and flow control set window size
- **connection-oriented:**
  - handshaking (exchange of control messages) initializes sender, receiver state before data exchange
- **flow controlled:**
  - sender will not overwhelm receiver

# TCP segment structure



# TCP sequence numbers, ACKs

## Sequence numbers:

- byte stream “number” of first byte in segment’s data

## Acknowledgements:


- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- A: TCP spec doesn’t say, - up to implementor

outgoing segment from sender

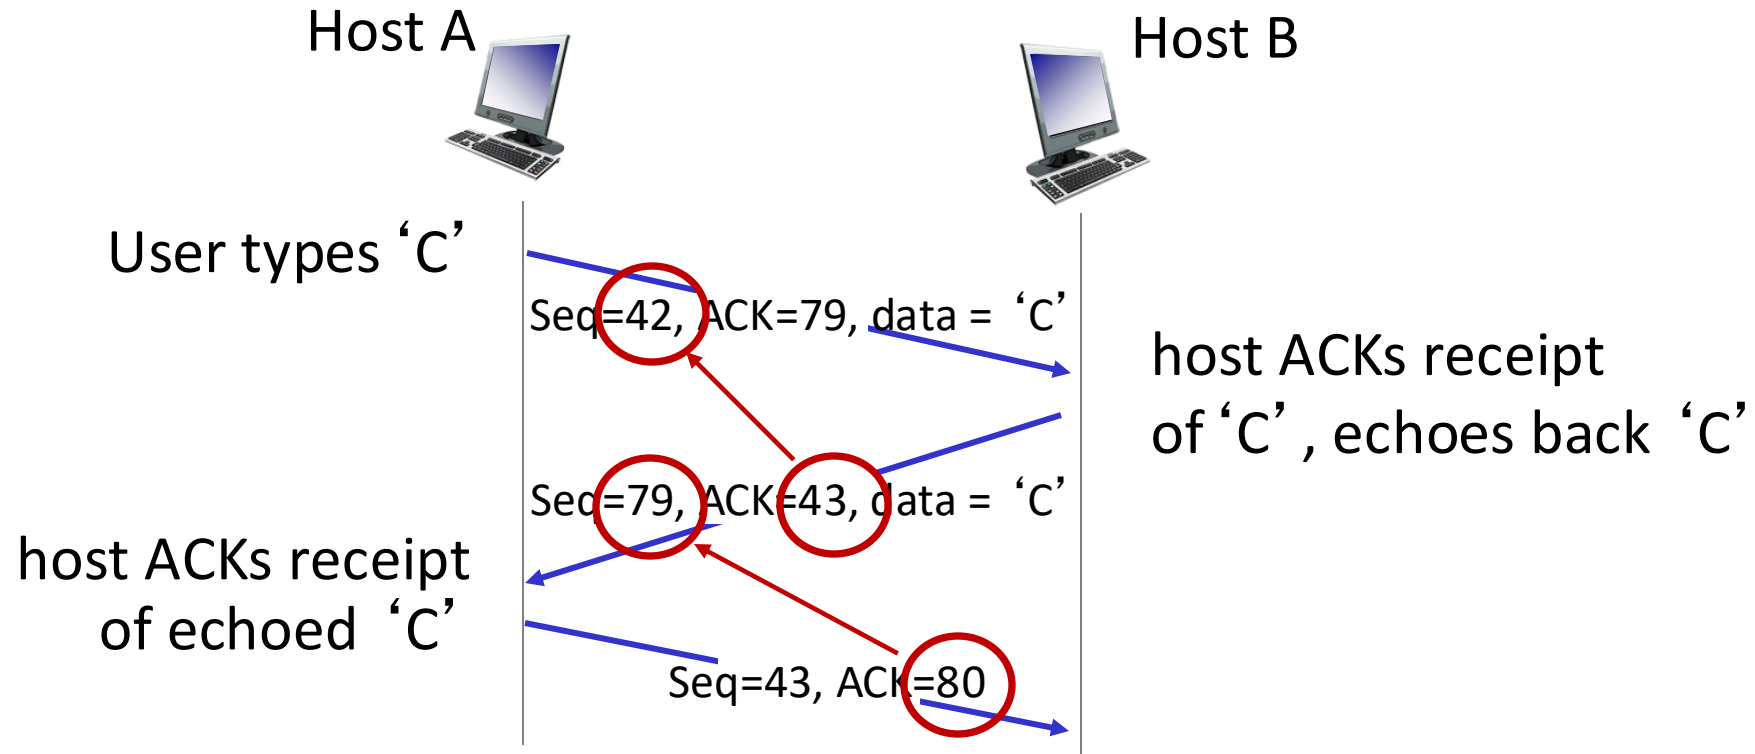
source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



outgoing segment from receiver

source port #	dest port #
sequence number	
acknowledgement number	
	<b>A</b> rwnd
checksum	urg pointer

# TCP sequence numbers, ACKs



simple telnet scenario

# TCP round trip time, timeout

Q: how to set TCP timeout value?

- longer than RTT, but RTT varies!
- *too short*: premature timeout, unnecessary retransmissions
- *too long*: slow reaction to segment loss

Q: how to estimate RTT?


- `SampleRTT`: measured time from segment transmission until ACK receipt
  - ignore retransmissions
- `SampleRTT` will vary, want estimated RTT “smoother”
  - average several *recent* measurements, not just current `SampleRTT`



# TCP round trip time, timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- exponential weighted moving average (EWMA)
- influence of past sample decreases exponentially fast
- typical value:  $\alpha = 0.125$



# TCP round trip time, timeout

- timeout interval: **EstimatedRTT** plus “safety margin”
  - large variation in **EstimatedRTT**: want a larger safety margin

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑  
estimated RTT

↑  
“safety margin”

- **DevRTT**: EWMA of **SampleRTT** deviation from **EstimatedRTT**:

$$\text{DevRTT} = (1 - \beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

# TCP Sender (simplified)

event: data received from application

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running
  - think of timer as for oldest unACKed segment
  - expiration interval: **TimeoutInterval**


event: timeout

- retransmit segment that caused timeout
- restart timer


event: ACK received

- if ACK acknowledges previously unACKed segments
  - update what is known to be ACKed
  - start timer if there are still unACKed segments


# TCP Receiver: ACK generation [RFC 5681]



# TCP: retransmission scenarios

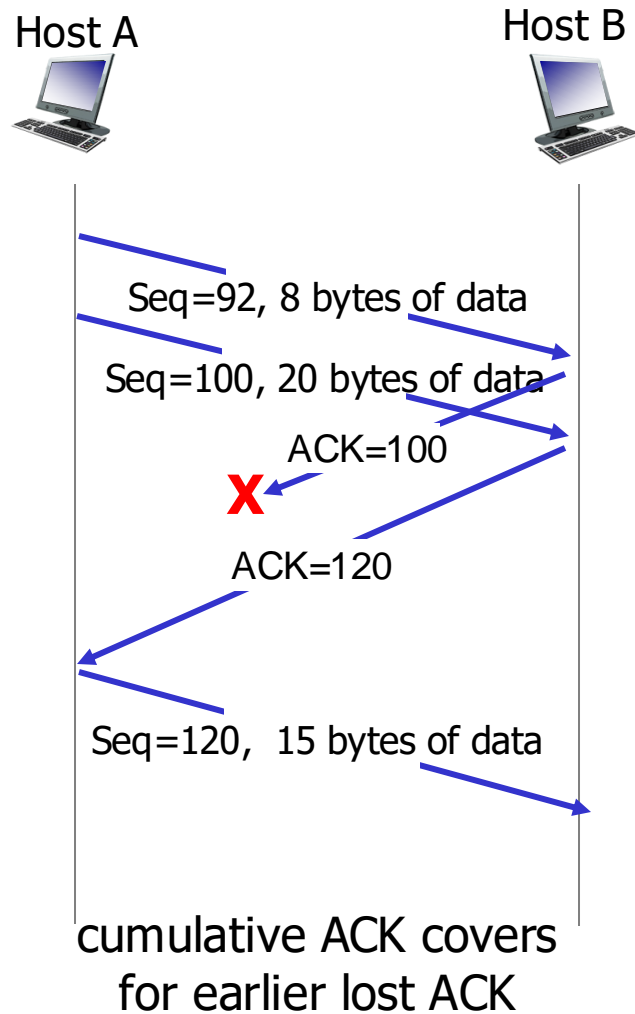


## lost ACK scenario



```
premature timeout
```

# TCP: retransmission scenarios



# TCP fast retransmit

## *TCP fast retransmit*

if sender receives 3 additional ACKs for same data (“triple duplicate ACKs”), resend unACKed segment with smallest seq #

- likely that unACKed segment lost, so don't wait for timeout



Receipt of three duplicate ACKs indicates 3 segments received after a missing segment – lost segment is likely. So retransmit!

